



COMPUTER SCIENCE

CAPSTONE REPORT - FALL 2022

A Project-based Assignment Package for an Undergraduate Operating Systems Course

Boyan Li

supervised by
Olivier Marin

Preface

I am a computer science student who also has an interest in using my skills in education. Operating Systems was a course that I found interesting but still complex and challenging to comprehend, and I am delighted to pursue this project as my capstone and contribute to the education of this course.

This paper is intended for educators teaching the Operating Systems course for undergraduate students. This paper presents a new teaching tool that reinforces students' understanding of the course material in an engaging way.

Acknowledgements

I would like to thank my professor, Olivier Marin, for assisting me through the development of this project, and my mother for supporting me through my difficulties. I would also like to thank the NYU Shanghai community as a whole, especially its faculty, Academic Affairs, Academic Resource Center, and Health & Wellness departments, for being helpful, thoughtful and forgiving throughout my undergraduate journey.

Abstract

Operating Systems is a course that is difficult to teach, as the concepts are opaque to most people. This paper presents a project-based assignment package to aid the teaching and learning of the course for undergraduate students. The assignment package requires the students to interact with operating system components by implementing a practical application, reinforcing their understanding of the course material during the process. The assignment package has a suitable complexity for undergraduate students and is engaging for students.

Keywords

Operating Systems, Project-based Learning

Contents

1	Introduction	5
2	Related Work	6
3	Solution	7
3.1	Overview	7
3.2	The Bitcoin data structure and related library functions	8
3.3	Part 0 - Warm up	10
3.4	Part 1 - Processes	12
3.5	Part 2, question 1 - Synchronization with semaphores and shared memory	14
3.6	Part 2, question 2 - Synchronization across multiple terminal windows	15
3.7	Part 3 - Threads	15
3.8	Part 4 - Files	17
4	Results and Discussion	18
4.1	Results	18
4.2	Evaluation	19
4.3	Discussion	20
5	Conclusion	21

1 Introduction

Operating systems are complex. An operating system is an architecture that comprises many components: scheduling, memory management, processes, interrupts, synchronization mechanisms, and so on. Each part is important, and yet opaque to a casual user. Therefore, teaching how an operating system works is a difficult task. Furthermore, Operating Systems is a fundamental course in many computer science programs, and knowledge and skills from the Operating Systems course are essential in many fields of study and work in computer science and related fields. As such, it is necessary to develop teaching approaches that allow students to not only understand operating system concepts but also practice the use of operating system functionalities in a practical scenario.

Many project-based approaches to teaching Operating Systems involve students writing code for components in an operating system kernel, or even writing an operating system from scratch. These approaches have several problems: their difficulty and workload can be overwhelming for an undergraduate student, and they may require the students to spend a lot of time debugging; while they allow students to gain hands-on experience with the inner workings of operating system components and algorithms, they do not offer an opportunity for students to experiment with the mechanics in a practical setting. In addition, many projects have highly connected components, and without a catch-up mechanism, a student who fails to complete a previous task may be unable to complete a later part.

This paper aims to present a project-based assignment package where students would complete an implementation of a practical application, namely, a Bitcoin mining simulator. The current implementation of the assignment package, written in C, contains 5 parts; the first part ensures students have a C compiling and running environment as well as basic C programming skills, while subsequent parts each cover a chapter of the Operating Systems course. For some parts of the assignment package, students will iteratively build off their solutions from previous parts, with a catch-up mechanism allowing them to continue with the assignments even if they did not finish, or performed poorly, on the prerequisite parts.

Although there are no experimental evaluations of the assignment package, this paper proposes a method and metrics to evaluate the effectiveness of the assignment package in practice.

An implementation of the assignment package, including template code and instructions for students, example solutions, as well as supplementary information and documentation, can be

accessed at <https://github.com/bl2437-nyu/os-assignment-package>.

2 Related Work

There have been various studies on teaching methods and teaching tools for the Operating Systems course.

[1] and [2] criticize old-fashioned lecture-based teaching methods, and both propose their improvements. Both studies have, among other proposals, a significant section involving practical or experimental teaching, suggesting the importance of hands-on exercises in Operating Systems teaching.

[3] proposes “problem-based learning”, and [4] proposes a “task-driven method”, both are methods centered around class discussions that improve students’ ability to analyze and solve problems. The authors of [3] state that their method is less efficient compared to traditional lectures, while [4] states that their method cannot be applied to all contents in the course material. The present study acknowledges the merits of said approach but believes that well-designed practical exercises also achieve the goal of improving problem-solving skills while being more effective and having better potential coverage of the course material.

The most common way to implement practical teaching is to have students write or modify operating system code, or write or modify a simulation of operating system features using a high-level language. [5] implements a simulation of operating system functionalities in Java; [6] implements an MOS (mini operating system) for use in an Operating Systems course, with the expectation that students have mastered Assembly language and C language before the course; [7] requires students to implement an operating system from scratch using Assembly and C. [7] notably tests their course with a group of students; in their test, most students are masters students, and some students express that “the hands-on labs are too difficult and too heavy”, implying their design was not suitable for undergraduate students. [2], however, included assignments that implement operating system components, as well as assignments that interact with operating system functionalities, in their proposed improvements to the teaching methods. The present study argues that, while implementing operating system components can allow students to gain a deeper understanding of the concepts, it can become overwhelmingly complex for an undergraduate student, especially when expecting the students to implement with Assembly language. On the other hand, practical tasks focusing on using operating system functionalities are

often more suitable for undergraduate students in terms of complexity, while offering insights on how operating system functionalities may be used in a practical application, which is much more useful for the students' future career.

Regarding tools for teaching Operation Systems courses, [8] explores the use of visualization tools to help students understand operation system algorithms; the present study acknowledges its merits but argues that practical teaching is equally effective at ensuring students' understanding of the course material. [9] proposes the use of virtualization software in the Operating Systems course, focusing on its ability to allow code to be demonstrated safely in a lecture. Considering the increasing accessibility of virtualization software over the years, the present study believes that virtualization software can also be applied in practical teaching, leveraging the safe and consistent environment that the virtual machine offers.

The present study aims to mitigate the disadvantages of the aforementioned works while integrating valuable insights from them with its proposed solution. More specifically, this paper presents a project-based teaching tool, in the form of an assignment package, focused on practical learning, with a suitable complexity such that an undergraduate student is able to understand the structure and logic of the project in its entirety.

3 Solution

The assignment package involves the students completing an implementation of a Bitcoin mining simulator, and by the end of the semester, the students would have a fully functional multi-threaded application.

3.1 Overview

The assignment package is split up into multiple parts. Part 0 is a warm-up part that confirms that students have basic C programming skills, and each subsequent part covers a section of the Operating Systems course.

A Bitcoin mining simulator is chosen for this assignment package for the following reasons:

- Bitcoin mining is, at its core, a relatively simple process to understand and implement. This allows the students to focus on designing and implementing functionalities required by the assignments without spending too much time understanding the Bitcoin mining logic.

- Bitcoin mining is a process that naturally justifies integrating the use of operating systems functionalities into the implementation, while still keeping the complexity at a manageable level. It is computationally expensive yet parallelizable, and requires minimal yet crucial inter-process and inter-thread communication, making it natural to integrate the use of processes, threads, and synchronization in a way that is meaningful but not overly complex. Other operating system functionalities can be integrated by introducing related processes surrounding Bitcoin mining, for example, writing a block or a blockchain to a file would make use of file I/O.
- Real Bitcoin mining applications also make use of various operating system functionalities. With an assignment package that implements a Bitcoin mining simulator, students are introduced to how operating system functionalities may be used in a real-world practical application.
- Bitcoin is a concept that students likely already know about, or wish to learn about, allowing the assignments to be relatable and possibly more engaging.

Students will be required to install a Linux virtual machine on which to compile and run all of the assignment code. The instructor should ideally provide the students with a suitable system image with C compiling tools pre-installed. Using a virtual machine ensures that every student has access to a C compiling and running environment, as well as ensuring that the behavior of the program is consistent across different devices. The installation and configuration of the virtual machine can also be seen as a problem-solving exercise for the students, albeit not explicitly related to the course material.

Students are assumed to have taken Data Structure and Computer Architecture courses and have basic C programming skills.

3.2 The Bitcoin data structure and related library functions

Throughout the assignment package, students will work with a Bitcoin data structure that is simplified but still captures the important components of a real Bitcoin block.

Various utility functions are also provided to students; these functions abstract many procedures related to Bitcoin and blockchain.

3.2.1 BitcoinBlock data structure

A `BitcoinBlock` struct used in this assignment package contains two main parts:

- A header, which has an identical format as the real Bitcoin’s block header. Crucially it contains a `nonce` which is modified during the “mining” process.
- A Merkle tree and a number of transactions.
 - For part 0 (“Warm up”), the tree is represented as a binary tree with nodes and pointers. This structure provides a reason for the students to use pointers and dynamic memory allocation during the assignment.
 - For subsequent parts, the tree is represented as a fixed-size array to fit a limited number of transactions. This avoids the use of pointers and makes it easier to store blocks in shared memory.

While real blockchains only link to each other by the header hash, this implementation also includes explicit information on where the previous and next block is stored. More specifically, each block in the blockchain is stored in a named shared memory, with its name containing the block’s header’s hash. Each block keeps the shared memory name for its previous and next blocks.

3.2.2 Project file structure and library functions

Header file	Source file	Description
<code>bitcoin_utils.h</code>	<code>bitcoin_utils.c</code>	Defines the Bitcoin block data structure, and includes functions related to algorithms used by Bitcoin as well as functions closely related to manipulating blocks and blockchains.
<code>data_utils.h</code>	<code>data_utils.c</code>	Contains utility functions related to manipulating data of, or related to, blocks and blockchains.
<code>debug_utils.h</code>	<code>debug_utils.c</code>	Contains functions for printing debug info.
<code>custom_errors.h</code>	<code>custom_errors.c</code>	Defines custom error codes that can be returned by certain functions, as well as a function that prints error messages, similar to <code>errno</code> and <code>perror()</code> in C standard library.
<code>sha2.h</code>	<code>sha2.c</code>	Enables computation of SHA2 hashes.
(none)	<code>main.c</code>	Entry point for the program. Students will write their code here.
(none)	<code>makefile</code>	Contain compile instructions for use by the command-line tool <code>make</code> .

Table 1: Description of files included in the code templates.

Table 1 shows a summary of files included in the template given to students for certain parts of the assignment package. Not all parts have a template: students are expected to modify their solutions for a previous part for some parts.

The various custom libraries provided implement many operations involving the Bitcoin blocks, transactions, and blockchains, as well as some debugging and error-handling functionalities. These help to abstract and simplify Bitcoin-related operations, allowing students to focus on knowledge and skills covered in the course material. All library functions are thoroughly documented, and have a manageable size; students are able to read through the source files to learn about the usage and implementations of these library functions.

3.2.3 Difficulty for block mining

Real Bitcoin uses a “difficulty” value to determine on average how many hashing attempts are needed to find a valid nonce; as of the time of writing this value is set to require around 2×10^{23} hashes. To keep the mining process practical for a regular home computer, a difficulty value is chosen that requires on average 1 million hashes to find a valid nonce, which usually takes a few seconds to compute on a typical laptop computer. Additional difficulty values with varying average required hashes are also defined, and students may choose a suitable value should they find the program running too fast or too slow.

3.3 Part 0 - Warm up

This warm-up section serves many purposes:

- Instruct students to set up a Linux virtual machine environment and C compiling environment, which would be used throughout the entire course.
- Make sure the students have basic skills for working with the C programming language.

These are of particular importance:

- Compiling and running a program using `gcc` and `make`
 - Manipulation of structs
 - Manipulation of pointers
 - Dynamic allocation of memory using `malloc()` and `free()`
- Introduce the students to the basics of the Bitcoin structure and Bitcoin mining.




Figure 1: A graphical representation of a 3-transaction Merkle tree, which students are instructed to construct as a part of this assignment.

3.3.1 Assignment breakdown

Students will be provided a template. The template contains code that creates a Bitcoin block and prints its header in a human-readable format.

Students will be asked to compile and run the code as-is, to confirm that they have correctly set up their environment and know how to compile and run a C program.

The students will then be instructed to perform the following tasks:

- Create a new Bitcoin block object.
- Create a Merkle tree, as illustrated in Figure 1, by creating the nodes and setting relevant pointers.
- Add transactions to the Merkle tree.
- Using provided utility functions, calculate and update this block's Merkle root.
- Write a loop to compute a valid nonce for the block by brute force (“mining”).

3.3.2 Rationale

For this part, the goal was to create a non-trivial task that covers basic skills of C programming. Since the rest of the assignment package is centered around Bitcoin mining, a warm-up task that performs basic Bitcoin operations is appropriate.

The students are instructed to compile and run the code as the first task after installing the virtual machine so that if any problems or difficulties arise, the student would likely have more

opportunity to troubleshoot the problem or find help.

Following that, the students are asked to write code that involves the creation and manipulation of custom struct types like `BitcoinBlock`, the creation and manipulation of pointers, and the use of `malloc()` and `free()`. This is because these skills are constantly being used throughout the rest of the assignment package, which means it is best to make sure that students have mastered the skills. The use of `malloc()` and `free()` would also give students a subtle introduction to memory management.

The students are then asked to write a brute force loop to mine the Bitcoin block. This brute force loop will appear in all subsequent parts of the assignment package, and students would need to modify the loops to insert various functionalities; it would be beneficial to ensure students' understanding of the basics of the process. At this point, the students also have an opportunity to gauge their device's computing power and choose a different difficulty value if needed.

Throughout this assignment, students are also encouraged to read source and header files for the provided libraries for definitions and documentation of custom types, values, and utility functions, as well as understand the compilation of a multi-file C program by examining the file structure and the `makefile` file.

3.4 Part 1 - Processes

This part covers the “processes” chapter of the Operating Systems course.

Students will be tested on the following skills:

- Creating child processes using `fork()`.
- Using the return value of `fork()` to distinguish between the parent process and the child processes.
- Using `getpid()` and `wait()`.

Considering semaphores and shared memories have not been introduced yet, the template includes code that allows the processes to communicate via signals, where upon receiving a `SIGUSR1` signal, a child process will exit, aborting the task that it is doing. Since signal-related APIs are not taught in the course, relevant code is mostly included in the template code, with an explanation of how to use `kill()` provided in the assignment instructions.




Figure 2: A graphical representation of the processes in part 1 and in part 2, question 1.

3.4.1 Assignment breakdown

Students will be given a template. The template code creates a random Bitcoin block and sets up signal handlers.

Students will be instructed to perform the following tasks:

- Using `fork()`, create several child processes.
- Write logic that distinguishes between the parent process and the child processes in an `if` statement.
- For the parent process, use `wait()` to wait for the first child process that completes its task and exits. Then, use `kill()` to send signals to the other processes, and wait for them to terminate as well.

Figure 2 shows a graphical representation of the relationship between processes created in this part.

3.4.2 Rationale

In the context of Bitcoin mining, once a process, thread, or perhaps another Bitcoin miner on the network, has successfully mined a block and attached it to the blockchain, other processes, threads, or miners should ideally stop and prepare to mine the next block. Thus, if we simply instructed the students to create child processes to run the brute force loop and exit when they find a valid nonce, it would be less optimal in the context. However, synchronization tools such as shared memory and semaphores have not yet been introduced to the students in this chapter.

This assignment package approaches this problem with a combination of `wait()` and signals. The parent process can `wait()` for a single process to exit (this process would be the one process that finishes first), and send a signal to the rest of the processes. This implementation would

require the students to record the PIDs returned from `fork()` and then compare them with the PID returned from `wait()`.

An earlier idea involved the parent process obtaining a nonce value from a child process via the exit code. However, only one byte of information can be obtained this way, which made this idea infeasible. Nevertheless, this idea might prove useful should a similar assignment package be designed with a different application.

While signals are an important part of an operating system, the use of signals and signal handlers in C programming is not covered in the Operating Systems course at NYU Shanghai. Thus, this assignment package provides students with all signal-related code within the template and introduces the use of `kill()` in the assignment instructions. Should this assignment package be used in an institution where the use of signals is covered in the curriculum, the assignment can easily be modified so that students need to write the code for signal handlers.

3.5 Part 2, question 1 - Synchronization with semaphores and shared memory

This part covers the “synchronization” chapter of the Operating Systems course.

Students will be tested on the following skills:

- Creating, referencing, and unlinking POSIX semaphores and shared memories.
- Using semaphores and shared memories to synchronize the processes.

3.5.1 Assignment breakdown

Students will be given a template. The template contains most of the code for a multi-process Bitcoin mining simulation. However, any synchronization logic has been left out of the template and requires the students to implement it.

Students will be instructed to implement synchronization between the parent and child processes. The program must be able to mine and attach multiple blocks without needing to terminate and re-create child processes.

3.5.2 Rationale

While testing early implementation ideas, it is discovered that it is too trivial to simply require all child processes to exit once one has reached a valid nonce - this requires a single boolean value in shared memory. As an attempt to make a non-trivial assignment that requires the use

of semaphores, this assignment requires the program to process multiple blocks without having the child processes exiting. This can be justified in the context of the application as avoiding the overhead of creating additional processes.

The example solution for this assignment uses two semaphore-based barriers to control execution order and one mutex to control write access to the shared memory. Students might be able to produce different implementations.

3.6 Part 2, question 2 - Synchronization across multiple terminal windows

This is a variation on part 2, question 1. Instead of multiple child processes of a single parent process running alongside each other, this question will have the student run multiple instances of the program from multiple terminal windows.

3.6.1 Assignment breakdown

Similar to Part 2, question 1, students will be given a template that contains most of the code but without synchronization logic. Students will be instructed to implement the synchronization logic between instances of the program.

3.6.2 Rationale

Despite not testing different skills compared to question 1, this question allows students to understand that named semaphores and named shared memories can be accessed by processes on the same operating system even if they do not share a common parent process. Students might also discover during the assignment that it is best practice to use a name for shared memories and semaphores that is unlikely to collide with names used by a different program.

Since it takes time for students to open multiple instances of the program in multiple terminal windows, and the program has no knowledge of the number of instances open, the template contains code that allows the student to type in the number of instances opened, and a barrier that releases all instances at once. Instructors may consider modifying this template code and instruct students to find a solution to this problem as a problem-solving challenge.

3.7 Part 3 - Threads

This part covers the “threads” chapter of the Operating systems course.

Students will be tested on the following skills:

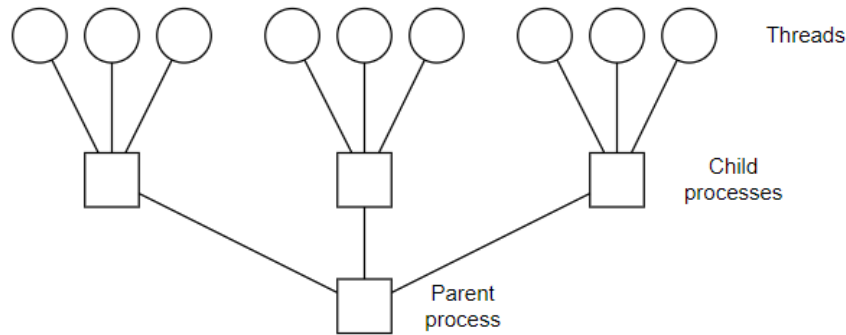


Figure 3: A graphical representation of the processes and threads in part 3.

- Creation of POSIX threads using `pthread_create()`
- Exiting a thread using `pthread_exit()`
- Waiting for threads to terminate using `pthread_join()`
- Synchronization of the parent process, child processes, and threads of child processes

While not required, students may also decide to use other POSIX threads APIs such as `pthread_mutex`.

3.7.1 Assignment breakdown

Unlike previous parts, students will complete this assignment by editing their solution from part 2, question 1. Students may request a copy of an example solution from part 2, question 1 as a template from their instructor, should they fail to complete that part.

Students will be instructed to update the child processes' code to create threads and have the threads perform the task of “mining” the Bitcoin blocks.

Figure 3 shows a graphical representation of the processes and threads created in this part.

3.7.2 Rationale

This is the first assignment where students are expected to iterate upon their previous solution. This allows students to experience the iterative development process, a process commonly employed in real-world development. However, students may decide to ask their instructor for a copy of the previous part's sample solution as a template, if they did not finish the previous assignment or were unsatisfied with their previous solution. It is at the discretion of the instructor whether to provide the sample solution.

Unlike part 2, question 1 where students were instructed not to terminate and re-create processes between blocks, this part does not enforce the same for threads. This is because the communication between threads from different processes is already a complex task. This can also be justified in the context of the application as threads require less overhead to create than processes.

3.8 Part 4 - Files

This part covers the “file management” chapter of the Operating Systems course.

Students will be tested on the following skills:

- Opening, closing, reading from, and writing to files, using `open()`, `close()`, `read()`, and `write()`.

3.8.1 Assignment breakdown

Students will complete this assignment by editing their solution from part 3. Students may request a copy of an example solution from part 3 as a template from their instructor, should they fail to complete that part.

Students will be instructed to update the code so that:

- When a block is mined and attached to the blockchain, write the blockchain’s data to a file. The write operation must not overwrite previously-written files, thus requiring a naming scheme to give each file written a unique name, and a metadata file containing the name of the most recently written file.
- When the program starts, try to read the metadata file, then the file with the name saved in the metadata, and load the blockchain data within into memory if it exists.

3.8.2 Rationale

Writing a blockchain to a file involves serialization and deserialization of binary data. While it could have been a good problem-solving challenge, asking students to design and implement a serialization scheme is rather complex and is ultimately a distraction from the course material. With that in mind, the reading and writing of a blockchain from or to a file have been packaged in two library functions, where students provide a file descriptor for the file involved. The students are still required to write code for reading from or writing to the metadata file.

The requirement for a naming scheme that provides unique names is also a minor problem-solving challenge, with potentially multiple different approaches that the students can take.

The use of a metadata file to store information about other files is also a practice that can be found in real-world applications.

4 Results and Discussion

4.1 Results

The assignment package proposed by this paper for the Operating Systems course has the following characteristics:

- **It emphasizes practical applications of operating system components.** Having students implement practical applications as assignments allows students to better understand how their course material may be applied in a real-world scenario, which is valuable for the student's future career. Existing approaches where students implement operating system algorithms or solve theoretical problems while allowing students to have a deeper understanding of the operating system components, fail to introduce the practical uses of the components.
- **Its difficulty and workload are suitable for an undergraduate student.** This assignment package uses the C programming language, a high-level language, which is much more manageable by an undergraduate student compared to assembly language. The library code is relatively small in size (approximately 1 000 lines of code, excluding `sha2.c` and `sha2.h`), allowing students to have a full understanding of the material they are provided with, in contrast to an operating system kernel which can have too much code for a student to understand, leaving a large portion of the material opaque to the student. Much of the Bitcoin-related operations are abstracted as library functions, reducing unnecessary complexity for the students.
- **It has a catch-up mechanism.** For parts where students are expected to modify their solutions from a previous part, if a student failed to complete the previous part, is unsatisfied with their previous solution, or otherwise finds difficulty modifying their previous code, they may request a template to modify for this part of the assignment. This can mitigate the cascading consequences of the student's previous mistakes.

- **It encourages problem-solving and creative thinking.** Students are required to solve various practical problems throughout the assignment package, whether related to the course material or not. Additionally, many tasks in the assignment package, as well as many library functionalities, have multiple different ways to approach; the students may be encouraged to compare the advantages and disadvantages of different designs and implementations.
- **It is relatable and engaging.** Throughout the course, the student would implement a fully functional Bitcoin mining simulator. Bitcoin and Bitcoin mining is a topic that students likely know about; students are likely to find the assignments relatable and the results rewarding.

4.2 Evaluation

Limited by the timeline of this project, an evaluation by experiments was not possible. If an experiment were to be conducted, the following methodology could be used:

A control group experiment would be conducted, where one class of Operating Systems students would receive traditional homework assignments for one semester as a control group, and another class of Operating Systems students would receive a modified set of assignments that includes this assignment package in the following semester as an experiment group.

A pre- and post-test could be used to evaluate the effectiveness of the assignment package in helping the students to understand the course material. A pre-test would be administered as a survey during or before the first lecture, or the instructor could verbally ask students about their knowledge of Operating Systems concepts during the first lecture. A post-test result would be taken from both the students' mid-term exam and final exam scores, as well as their performance in the assignment submissions. The pre- and post-test results would be compared to assess students' performance under their respective conditions.

In addition, qualitative feedback would be collected from the students. Surveys and interviews would be conducted, and students would be encouraged to submit written feedback alongside their assignment submissions. This feedback would be used to assess the quality of the assignments and student engagement.

Some of the questions that could be asked during a survey or interview, or be instructed students to write in feedback alongside their assignment submissions, could include:

- How difficult was the assignment? What was/were the main difficulty(-ies)?

- Were the assignment prompt and code template clear and easy to understand? Were they helpful?
- Did you try to read the code in the provided libraries? How well did you understand it?
- Are the assignments, or certain parts of them, interesting? Are there parts that made you consider factors beyond the requirements of the assignments?
- Did the assignments help you understand, or help you reinforce your knowledge of, the course material?

4.3 Discussion

Despite the advantages of the assignment package, there are significant limitations that need to be addressed:

- **It does not cover all the chapters in the course.** The current implementation of the assignment package covers processes, synchronization (using semaphores and shared memories), threads, and files. This still leaves significant portions of the course unaccounted for, such as scheduling, memory management, and networking. While certain chapters can be more easily integrated into the assignment package, others might be more difficult to incorporate. The following are potential avenues for expansion that are yet to be implemented:
 - Networking can be incorporated by requiring the program to communicate with other programs to simulate a Bitcoin network. The implementation may be decentralized like the real Bitcoin, or may be modified to use a server-client model for simplicity.
 - Memory management can be simulated by arranging transaction data in a byte array in the Bitcoin block structure. Instead of using memory locations or C pointers, indexes into the byte array would be used to refer to a location in the array instead.
- **It does not cover edge-case scenarios and does not cover all the details in the course material.** Implementation of a practical application is likely to take the least resistant path, which can leave part of the course material unaccounted for in the assignments. For example, when implementing reading and writing of files, a student is unlikely to implement in a way where multiple processes write to the same file at the same time, nor is it likely to use `dup()` to duplicate a file descriptor. However, these edge-case scenarios and functionalities can still be important for students to understand. As such, this assignment

package should not be used on its own, instead, it should be used in conjunction with other practice questions, lab exercises, or assignments that account for the missing parts. It may be possible to improve the design of the assignment package to improve the coverage of these details.

A minor challenge during the design and implementation process of this assignment package is to navigate the various requirements, goals, and technical limitations, all of which place limitations on how the package can be implemented. Some of the requirements are that the assignments have to cover the course material while still feeling natural within the context of the Bitcoin mining simulator; they have to be challenging but not overwhelming; the library code has to be relatively robust. Some technical limitations include the inability to use pointers inside shared memory, which also conflicts with the warm-up part's requirement to incorporate pointers. While none of these caused severe difficulties during the design and implementation process, they did require additional consideration, and sometimes multiple iterations of different implementations; notably, the Bitcoin block structure has undergone 5 iterations so far.

During the course of this project, the author has been made aware of alternative applications to theme this assignment package around. Some of these applications include a file system simulator, a multi-threaded web server, or a database server. However, each of these applications has properties that make them unsuitable for this assignment package, for example, a file system simulator is much more limited in scope and does not cover all chapters of the course material; a web server requires parts of the assignment to be in a different order as the course material, with networking required near the beginning, while also being less compatible with an iterative development process that this paper aims to implement. The author believes that a Bitcoin mining simulator remains an optimal selection for this assignment package.

5 Conclusion

This paper proposes a project-based assignment package for an undergraduate Operating Systems course, where students implement a Bitcoin mining simulator using various operating system concepts. The assignment package allows students to gain an understanding of how operating system functionalities may be applied in a practical scenario, with a workload and difficulty suitable for undergraduate students. The assignment package addresses a gap in project-based teaching tools suitable for undergraduate students.

The current implementation contains 5 parts, covering 4 chapters of the Operating Systems course. Code templates and custom libraries are provided to students to manage the complexity of the assignments. Some parts of the assignments require students to incrementally iterate on their previous solutions, and a catch-up mechanism is provided for students that fail to complete a previous assignment.

While it does not currently cover all chapters of the course nor does it cover every detail in a chapter, it has potential to be expanded and improved. An experiment to evaluate the assignment package's effectiveness in a teaching environment was not performed, but an appropriate methodology for evaluation was presented.

References

- [1] W. Qingqiang and C. Langcai, "Teaching mode of operating system course for undergraduates majoring in computer sciences," in *2009 4th International Conference on Computer Science & Education*, 2009, pp. 1412–1415.
- [2] J. Niu, Z. Hu, and X. Xiao, "Study on curriculum reform of computer operating system," in *2009 First International Workshop on Education Technology and Computer Science*, vol. 1, 2009, pp. 262–264.
- [3] H. Yi-Ran, Z. Cheng, Y. Feng, and Y. Meng-Xiao, "Research on teaching operating systems course using problem-based learning," in *2010 5th International Conference on Computer Science & Education*, 2010, pp. 691–694.
- [4] L. Fang and X. Li, "Research application of task-driven method in teaching of the course operating system," in *2009 International Conference on Education Technology and Computer*, 2009, pp. 164–166.
- [5] J. Oh and D. Mosse, "Doritos (distributed object-based real-time instructional operating system): a complete package for teaching principles and practices of real-time operating system," in *FIE '98. 28th Annual Frontiers in Education Conference. Moving from 'Teacher-Centered' to 'Learner-Centered' Education. Conference Proceedings (Cat. No.98CH36214)*, vol. 1, 1998, pp. 313 vol.1–.
- [6] H. Li, C. Yin, Y. Xu, and Q. Guo, "Construction of the practical teaching system on operating systems course," in *2010 Second International Workshop on Education Technology and Computer Science*, vol. 1, 2010, pp. 405–408.
- [7] S.-L. Tsao, "A practical implementation course of operating systems: Curriculum design and teaching experiences," in *2008 14th IEEE International Conference on Parallel and Distributed Systems*, 2008, pp. 768–772.
- [8] A. Alharbi, F. Henskens, and M. Hannaford, "Integrated standard environment for the teaching and learning of operating systems algorithms using visualizations," in *2010 Fifth International Multi-conference on Computing in the Global Information Technology*, 2010, pp. 205–208.
- [9] D. Dobrilovic and Z. Stojanov, "Using virtualization software in operating systems course," in *2006 International Conference on Information Technology: Research and Education*, 2006, pp. 222–226.